TECHNISCHE UNIVERSITÄT
KAISERSLAUTERN

Building an LLVM backend for TriCore architecture

Kumail Ahmed

Supervisors:
Prof. Dr. Wolfgang Kunz & M.Sc Ammar Ben Khadra

This thesis is presented as part of the requirements for the conferral of the degree:

M.Sc Electrical and Computer Engineering

University of Kaiserslautern
Department of Electrical and Computer Engineering

March 2016

Declaration

I, Kumail Ahmed, declare that this thesis submitted in fulfilment of the requirements for the conferral of the degree M.Sc Electrical and Computer Engineering, from the University of Kaiserslautern, is wholly my own work unless otherwise referenced or acknowledged. This document has not been submitted for qualifications at any other academic institution.

—————————————————————

Kumail Ahmed
April 20, 2016

Contents

Chapter 1

Introduction

The computational power of processors is doubling every eighteen months. The expanse of embedded systems spreads over all domains of applications - aerospace, medical, automotive, etc. With the ever increasing complexity of hardware and software, the focus on speed, performance is becoming more important. This has lead to development of new processor architectures that can cope with these ever-growing needs.

Compilers play an integral role at the point where the hardware meets the software. They have become a well established research domain for hardware research. Essentially, the job of a compiler is to convert a high-level programming source code into a target language in a reliable fashion. Compiler architecture is divided into three parts:

1. Front End : The job of the front end is to analyse the structure of the high-level source code and build an intermediate representation (IR) of the code. Checks for syntax and semantic errors are also performed in this phase.

2. Middle End : Performs optimizations on the intermediate representation.

3. Back End : The backend is responsible for generating architecture specific code from the intermediate representation provided by the middle end.

1.1 Thesis aim

The goal of this thesis is to develop a backend for a TriCore architecture using Low Level Virtual Machine (LLVM 3.7). LLVM provides a modular compiler infrastructure that provides this frontend/backend interface.

As mentioned before, the responsibility of the backend is to convert the target agnostic IR representation into system-dependent representation, and generate assembly code as a result. This conversion requires a lot of features that are mostly

hidden from an application programmer. Some of these features include calling convention layout, memory layout, register allocation, instruction selection.

1.2 Thesis outline

This thesis is divided into five chapters. The first chapter is an introductory overview of the work. The second chapter gives an introduction about the TriCore architecture. Chapter 3 gives an introduction to the LLVM compiler infrastructure. Chapter 4 provides a description of the backend implementation. The thesis concludes with a conclusion and discusses some future works.

Chapter 2

TriCore Architecture

Infineon started the first generation of Tricore microprocessor in 1999 under the trademark AUDO (AUtomotive Unified-ProcessOr). Since 1999, the company has advanced the TriCore technology, and currently the 4th generation TriCore chip is sold under the trademark of AUDO MAX. Currently, Tricore is the only single-core 32-bit architecture that is optimized for real-time embedded systems. TriCore unifies real-time responsiveness, computational power of a DSP, and high performance implementation of the RISC load-store architecture into a single core.

TriCore provides simplified instruction fetching as the entire architecture is represented in a 32-bit instruction format. In addition to these 32-bit instructions, there are also 16-bit instructions for more frequent usage. These instruction can be used to reduce code size, memory overhead, system requirement, and power cost.

The real-time capability of the TriCore is defined by the fast context switching time and low latency. The interrupt latency is minimized by avoiding long multi-cycle instructions.This makes TriCore a wise choice for in a real-time application. TriCore also contains multiply-accumulate units that speed up DSP calculations.

This chapter describes the key components of the TriCore ISA that are essential in the understanding of the backend design.

2.1 Register description

Tricore consists of following registers:

1. 32 General Purpose Registers (GPRs)

2. Program Counter (PC)

3. Previous Context Information Register (PCXI)

4. Program Status Word (PSW)

The PC, PCXI, and PSW registers play an important role in storing and restoring of task context[1].

The 32 GPRs are divided into two types, i.e the so-called Address registers and Data registers. The Address registers are used for pointer arthimatics, while data registers are used for integral/floating type calculation. This peculiar Address/Data distinction creates a problem that would be discussed in chapter 4 in the section of calling convention implementation. The following table shows the registers and their special functions:

Data Registers	Address Registers	System Registers
D15 (Implicit Data)	A15 (Implicit Base Address)	PC
D14	A14	PCXI
D13	A13	PSW
D12	A12	
D11	A11 (Return Address)	
D10	A10 (Stack Return)	
D9	A9 (Global Address Register)	
D8	A8 (Global Address Register)	
D7	A7	
D6	A5	
D5	A5	
D4	A4	
D3	A3	
D2	A2	
D1	A1 (Global Address Register)	
D0	A0 (Global Address Register)	

Table 2.1: TriCore registers

D15 and A15 are the implicit registers that are normally used by 16-bit instructions. The registers A0, A1, A8, and A9 are designated as global registers, and they are neither saved nor restored between function calls. By convention A0 and A1 register are reserved for compiler use and A8 and A9 are reserved for application usages. A11 holds the return address from jump and call instructions.

Finally, the two distinct colors show the two respective task context. Register A10-A15, D8-D15, PSW, and PCXI belong to the upper context, while registers A2-A7, and D0-D7 belong to the lower context. The upper context is automatically restored using the RET instruction. The lower context is not preserved automatically[2].

Moreover these GPRs can also combine in an "odd-even" pair to form a 64-bit register. There are no intrinsic real 64-bit registers in TriCore, hence for performing calculations that require 64-bit manipulation, an "extended register" is created by the "odd-even" combination. E0 is defined as [D1-D0], E2 is defined as [D3-D2], and so on. By convention, extended registers for the address type are named as P[0], P[2],

and so on. Extended registers are used when multiplying large numbers or passing 64-bit arguments as a formal argument. More about this would be discussed in the calling convention section.

2.2 Supported data types

The Tricore Instruction set supports the following data types:

1. Boolean : mostly used in conditonal jumps and logical instructions.

2. Bit String : produced using logical, and shift instrucitons.

3. Byte : an 8-bit value

4. Signed Fraction : comes in three varients 16-bit, 32-bit, and 64-bit. Mostly used in DSP instructions.

5. Address : a pointer value.

6. Signed and unsigned integers : a 32-bit value that can either be zero- or sign-extended. short signed and unsinged integers are sign-extended or zero-extended when loaded from memory to a register.

7. IEEE-754 Single precision Floating-point number

Hardare support for IEEE-754 floating point numbers and long long integers in provided wit the basic TriCore ISA. Hence, a specific coprocessor implementation is required in order to extend the ISA.

The address is always a 32-bit unsigned value that points to a memory address. In C parlance, it is simply a pointer variable. Hence, if the following code is executed, the ptr variable would always be stored in an address type register, and always be a positive integer.

```
int a = 10;
int *ptr = &a; // ptr holds the address of variable a
```

Listing 2.1: Pointer example

2.3 Data formats, and alignment

The 32-bit TriCore registers can be loaded as a byte, a half-word, a word or as a double-word. The particular load/store instructions define whether the value is loaded/stored as a sign extended number or a zero extended number. For example,

calling LD.W loads a word with sign extension and LD.WU loads a word with zero extension.

Figure 2.1: Supported data formats in TriCore. Adapted from [2]

Tricore uses little endian byte ordering. Byte ordering is performed either at 1,2 or 4 byte boundaries. The following table lists the alignment information for the primitive data types in C.

Data Type	Size	Alignment
char	1	1
short	2	2
int	4	4
long	4	4
long long	8	4
float	4	4

Table 2.2: Alignment information for primitive data types

2.4 Instruction Set

TriCore instruction set is divided into two formats: 16-bit opcode format and 32-bit opcode format. These formats are then further divided into different types depending on the type of data type modifiers. In total, there are 12 sub-formats for the 16-bit opcodes and 25 sub-formats for the 32-bit instructions. These two type formats can be distinguished on the basis of the zeroth bit. The zeroth bit for a 16-bit format is always zero, and always one for 32-bit format.

For example, one of the frequent occurring sub-format for a 32-bit opcode is called RC. The letter "R" means register, and "C" means a constant. Hence instructions such as a register and constant ADD would use an instruction that is of

the type RC. Similarly, a register-register ADD would use a type RR sub-format. This distinction allows grouping common instructions together when writing them in the LLVM backend.

The 16-bit opcode format mostly supplements common instructions in the ISA. for example, eight variants exist for the ADD instruction in TriCore. These variants are characterised according to the range of constant value they can take as an input, the type of register they point to, and the opcode format. Consider the following example,

```
int a = 10;
a = a + 10;
```

Listing 2.2: Addition example

Assuming that constant elimination is turned off, the backend can either generate a 32-bit instruction or a 16-bit instruction. Furthermore, it can also chose whether to put the result in an implicit register (i.e. D15) or any other data register. Considering these nuances would allow in the generation of a more efficient code in terms of size and speed.

TriCore provides several addressing mode for loading and storing data to and from memory. The most commonly used mode are :

1. Base + Short Offset (Format BO)

2. Base + Long Offset (Format BOL)

3. Pre/Post Increment (Format BO)

The addressing modes define the effective address for a load/store instruction, and then update the value of the base pointer. In this thesis, the first two modes are mostly used. Just like in the case of ADD instruction, Load and store instructions can also be divided into several types. Furthermore, TriCore defines a whole range of different loads and store depending on the primitive data type.

Opcode format for a jump instruction is characterised with an initial "B". Hence a jump that compares a register and a constant would be of the type BRC. On the same line, a register-register compare is of the type BRR. Generic jumps are in 32-bit opcode format, while jumps that compare with zero comes in 16-bit opcode format. Comparing to other RISC architectures like ARM, the jump instruction in TriCore performs comparison and jump in the same instruction. Hence, there is no need to use the common status register. In most cases the result for a jump instruction in written in the implicit register, but instruction also exist of for writing them into a register of a different type.

TriCore offers only a single shifting instruction For Arthimatic and logical shifting. The performing a left shift a positive shift value is given, while for a right shift, the value is passed in negative. The shift counter is a 6-bit signed number, hence allowing values from -32 to +31, allowing 31 bits shift to the left and 32 bits shift to the right.

Chapter 3

LLVM Structure

LLVM is a "collection of modular and reusable compiler and tool-chain technologies"[3]. More specifically, it a compiler infra-structure that provides a front-end (parser and lexer), optimizations, and a back-end (that converts the IR representation into machine code). LLVM also provides a range of optimizations including compile-time, link-time, and run-time optimizations. The project started at the university of Illinois by Chris Lattner, and has now grown into a huge open-source project including a vast range of tool-chain technologies such as JIT systems, debuggers, optimizers, interpreter, etc.

3.1   LLVM Design

A brief introduction of compiler design was given in chapter 1. A classical compiler is divided into a front-end, optimizer (middle-end), and a back-end. In the context of LLVM, a significant benefit of such a design is retargetability. LLVM uses the same optimizer for any arbitrary high-level language that has to be translated into machine code.

This technique considerably reduces the amount of code that has to be written when writing a new compiler from scratch. Hence, without retargetability, a compiler that support M programming languages, and N back-ends, M*N compilers ought to be written. LLVM reduces this job to M+N.

The three-phase design also offers encapsulation between the different parts of the compiler. This leads to another major benefit that different open-source communities can focus on specific parts of the tool-chain. The expertise required for designing a front-end are different from implementing a backend. This allows in the development of efficient optimizations in a smaller amount of time. In the case of proprietary compilers the efficiency is defined by the amount of budget that a company invests and the target market of the compiler designer. An example of a proprietary compiler is HP aC++.

Figure 3.1: LLVM retargetability

3.1.1 Frontend

LLVM was initially written as a code generator for GCC. As the project evolved, new frontends starts to emerge. Currently, LLVM supports a wide variety of high-level languages including C, C++, Pascal, Ada, Fortan, Objective-C, and Swift. The most poupular frontend used with LLVM is called clang. It is also developed as an opensource project and support byte code generation for C, C++, Objective-C, and Objectice-C++. clang has the following advantages:

1. More readable diagnostic messages

2. Compatible with GCC, and lower memory usage.

3. Supports a wide variety of clients including services like static code analysis, and code refactoring.

   Another known frontend was llvm-gcc. It was a modified version of the original GCC compiler that uses LLVM as its backend and had support for C and Onjective-C language. This project is now deprecated with DragonEgg that is attached with GCC 4.5 and above as plugin. The DragonEgg extension also support languages like Ada and Fortran, although results are poor for Ada on newer GCC verison. Compared to clang, the diagnostic information from DragonEgg is poor, and only compiles a "reasonable amount" of Objectice-C++ code.

3.1.2 Backend

LLVM 3.7 has code generation support for several architectures including AArch64, AMDGPU, ARM, BGP, Hexagon, MSP430, Mips, NVPTX, PowerPC, Space, SystemZ, and XCore. The latest release also has support for AVR, but it is in an experimental stage. Backend targets can be accessed from $llvm/lib/Target$ folder

3.2 LLVM Intermediate Representation

The most valuable aspect of LLVM is its intermediate representation(IR). LLVM intermediate representation is a type of human-readable assembly language yet powerful. It is designed to be a generic IR that on one hand can represent low-level machine information, and on the other hand encapsulate high-level language ideas. While the front-end and back-end are restricted by the constraints of the source language and target architecture respectively, an IR provides complete freedom to the optimizer both in terms of expressiveness and performance.

```
define void @foo(i32 %d) #0 {
entry:
  %d.addr = alloca i32, align 4
  %a = alloca i32, align 4
  %b = alloca i32, align 4
  %c = alloca i32, align 4
  store i32 %d, i32* %d.addr, align 4
  store i32 230, i32* %a, align 4
  store i32 210, i32* %b, align 4
  %0 = load i32, i32* %a, align 4
  %1 = load i32, i32* %d.addr, align 4
  %add = add nsw i32 %0, %1
  store i32 %add, i32* %c, align 4
  ret void
}


// C code for this IR
void foo(int d) {
  int a=230,b=210,c;
  c = a + d;
}
```

Listing 3.1: LLVM intermediate representation example

As can be seen from the listing, LLVM IR is a similar to a RISC instruction set. It had instruction for loading and storing from and to the the memory. All other operations are performed in virtual registers (%0, %1, %add). LLVM has a static single assignment (SSA) form, hence it requires that each virtual register is assigned exactly once. All instruction are in the three address form, where the left-hand operand defines the destination register, while on the right-hand side there are exactly two operands.

3.3   Target-Independent Code Generator

The target-independent code generator contains reusable LLVM components that provides the framework for generating machine code for specific target platforms. There are six major components of this framework:

1. **Abstract target description** defines a target at an abstract level. It contain information about the instructions and registers that a architecture might support.

2. **Classes for code generation** are abstract classes implemented in */llvm/lib/CodeGen*. These are abstract classes that must be implemented for a specific target architecture.

3. **Classes for MC layer** define textual information that is inserted into a generated assembly. This information include block labels, constrain information, header and footer for assembly, etc.

4. **Classes for target-independent code-generation algorithms** define the different phases of the backend including register allocation, instruction selection, stack representation, etc.

5. Implementation of abstract target description for a specific target. Any concrete architecture (i.e. TriCore) must inherit from this abstract description.

6. Target independent implementation of Just-in-Time compiler (optional).

In the development of this thesis work, the MC layer classes were implemented both for assembly generation and object file generation. The exact details of the implementation are presented in the next chapter.

3.4   TableGen tool

Before writing a backend, the first prerequisite is to extract the ISA and ABI information present in TriCore documentation and write it for LLVM. The mechanism that LLVM uses to express this information is written into target description files. For writing this target description files LLVM uses a tool called TableGen.

TableGen is a domain specific language that is used for developing human readable record for instruction and register description files for a specific target. Tblgen greatly reduces the amount of work that has to be done in describing instructions by allowing code reuse. A backend engineer will describe the target that is to be implemented in .td files that will be translated into C++ code by LLVM. Henceforth,

generated C++ can be used in describing the overall process of instruction selection and register allocation. Secondly, description files allows code maintainability and relaiablity due to its modular design.

Although TableGen has a defined semantic, not much information is available apart from [4]. The best source to undersrtand its functionality is to read implementation performed in other target architectures.

TableGen operates by performing two distinct passes. A pass in llvm is defined as a piece of code that performs optimizations and transformations. The first pass is an expansion of a the templates defined by the back-end engineer. This expansion pass is target-independent, and is predominantly a huge switch-case structure inside the *build/lib/Target/TriCore* folder in the LLVM hierarchy. The second pass adds target dependent information and is packaged inside various .inc files for a specific backend.

Figure 3.2: Class hierarchy for the table generator file systems. As a convention the name of the file starts with the name of the architecture. The files are grouped together according to their common functionality. In the fist pass, tblgen expands the macros. In the second pass target-specific information is incorporated with these exapnded marcos.

Chapter 4

Backend Design

In essence, the backend constitutes a number of passes that converts the LLVM intermediate representation into machine code. As mentioned in chapter 3, every backend has to inherit from a common set of abstract classes to perform target-specific code emission. A thorough discussion is given in [5]. At first, LLVM IR is passed to the backend. As the code generation process progresses forward, the IR is converted into TriCore-specific representation. In toto, there are seven steps that are performed by LLVM. The diagram below gives an overview of the steps to output either an assembly or an object file as a final product.
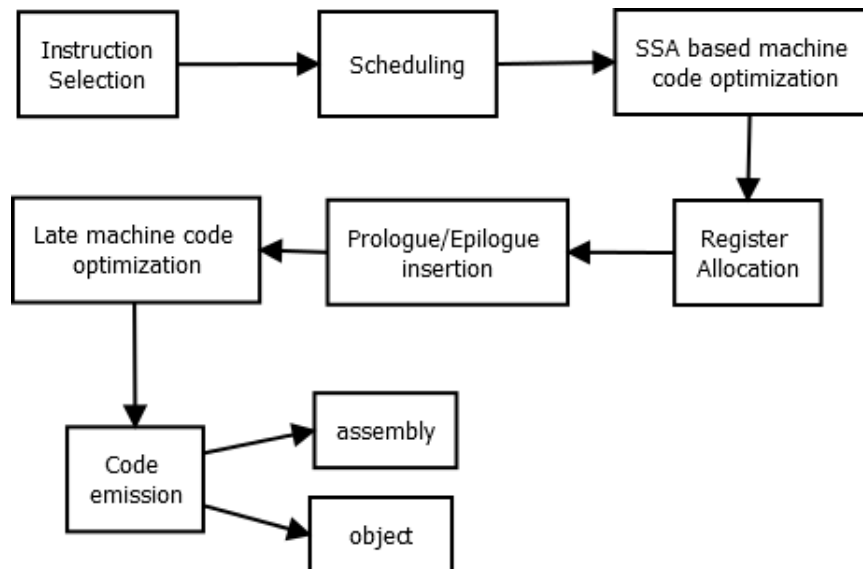

Figure 4.1: Example of a DAG

Each phase is called a pass in LLVM. The phases that are marked as *optimization* are crucial for increasing the code quality, while other passes are called *superpasses* that are made of up many smaller passes defined inside the LLVM framework. These superpasses are the most essential to be implemented in order to get a working backend.

In this chapter, a detailed overview is given for the above mentioned steps and related information about the essential prerequisites.

4.1 Tblgen Files

4.1.1 Instruction Pattern

Before we dive into the LLVM's code generation process, it is necessary to understand how an instruction is described in tblgen. Tblgen essentially attaches a set of input and output DAGs (Direct Acyclic Graph) either to a LLVM defined pattern or a user-defined pattern. A DAG is a collection of vertexes and edges where an end node Z can never loop back to a start node A.

Figure 4.2: Example of a DAG

Similarly, tblgen uses DAGs to represent an instruction. For example, consider the pattern selection defined below for an ADD instruction in TriCore. It contains two different DAG, i.e. an output DAG defined by the keyword *outs*, and an input DAG defined by the *ins* keyword. Next the assembly string is supplied, that is printed once at the code emission stage. Lastly comes in the matching pattern defined inside the [ ] brackets. The pattern defined below assigns the output of the ADD node to the register $d.

```
def ADDrr : RR<0x0B, 0x00, (outs DataRegs:$d),
        (ins DataRegs:$s1, DataRegs:$s2),
        "add $d, $s1, $s2",
        [(set i32:$d, (add i32:$s1, i32:$s2))]>;
```

Listing 4.1: Pattern selection example

4.1.2 Register Classes

The ADD instruction mentioned above also defines a register class that it inherits from. TriCore register information is present in register classes. The register class defines the data type, name, and optional dwarf information for a given register. As mentioned in chapter 2, TriCore has two different types of registers, i.e. the address registers and the data registers. Each type of register inherits from its respective class as described in the hierarchy in figure 4.3.

TriCoreRegWithSubregs class is used to define registers that hold 64 bit data values. As there is no true 64.bit register available in TriCore, two 32-bit registers have to be combined together. The process of doing this in tblgen is governed by the idea of sub-registers. A sub-register can either be the lower byte or the higher byte in the context of 64 bits. These sub-registers are later assigned to an extended registers by assigning individual sub-registers to an extended register.

```
let SubRegIndices = [subreg_even, subreg_odd] in {
   def E0 : TriCoreRegWithSubregs<0, "e0", [D0,D1] >, DwarfRegNum<[32]>;
   }
```

Listing 4.2: 64-bit register definition

In the example above, Extended register E0 is made up of two sub-registers that called even and odd. The odd register in this case is D1, while the even register is D0. Furthermore, D0 is the MSB (Most Significant Byte) while D1 defines the LSB (Least Significant Byte) for TriCore. The DwarfRegNum comes from the TriCore EABI information giving a unique debug code to each register.
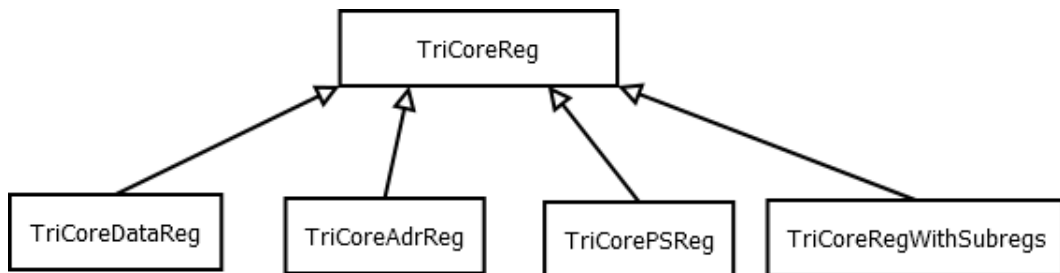
Figure 4.3: Register class hierarchy

4.1.3 Instruction Format

Continuing with the ADD instruction example above, The initial $RR <$ represents the instruction format this instruction belongs to. An instruction format allows the representation of appropriate byte codes for a specific instruction, the output and input DAGS, the assembly string, and the pattern that has to be matched. For example consider an AND instruction.

```
Instruction Description
AND     D[c], D[a], const9

Operation in RTL Format
D[c] = D[a] & zero_extended(const9)
```

Figure 4.4: Instruction format for AND instruction

An instruction format is a direct representation of TriCore instruction type in Tblgen language. This AND instruction takes in a data registers and a constant value and outputs the result into a third register. Furthermore, it defines two opcodes that have to assigned in order to perform a bit-wise AND operation. An example in tblgen for such an instruction format is as follows:

```
class RC<bits<8> op1, bits<7> op2, dag outs, dag ins, string asmstr,
             list<dag> pattern> : T32<outs, ins, asmstr, pattern> {
  bits<4> s1;
  bits<4> d;
  bits<9> const9;

  let Inst{7-0} = op1;
  let Inst{11-8} = s1;
  let Inst{20-12} = const9;
  let Inst{27-21} = op2;
  let Inst{31-28} = d;
}
```

Listing 4.3: Instruction format for a RC-type instruction

Defining a class for each specific instruction format provides encapsulation and inheritance capabilities. This allows in writing more legible and efficient code for other instructions that inherit similar properties.

4.1.4 Representing Constants

Many instructions in TriCore take in constants as parameters. A constant is either a signed or a zero type value. A singed number, as its name suggest, can have both negative and positive value, while a zero type integer holds any value greater or equal to zero. In Tblgen parlance, these properties are supplied using constraints. The instruction described in figure 4.4 takes a 9-bit constant as a parameter. Such an operator constraint is written as:

```
def immZExt9 : ImmLeaf<i32, [{return Imm == (Imm & 0x1ff);}]>;
def immSExt4 : PatLeaf<(imm), [{ return isInt<4>(N->getSExtValue()); }]>;
```

Listing 4.4: constant constraints

A 9-bit zero constant can have a value $0 \leq Imm \leq 2^9$. For a signed constant, the most significant byte is used to describe the sign. Hence, the range for a 4-bit number is $-2^3 \leq Imm \leq 2^3 - 1$.

4.1.5 Summary

An overall hierarchy of tblgen files is present in figure 3.2. The organization of tblgen files is as follows

- **TriCoreInstrInfo.td**,
  TriCoreInstrFormats.td: Contains information related to TriCore instructions and their formats

- **TriCoreOperators.td**: Contains information related to TriCore addressing and constant operator, parameter constraints, and user-defined TriCore nodes.

- **TriCoreRegisterInfo.td**: Contains information related to TriCore register description including type information, alignment, and debug information.

- **TriCoreCallingConv.td**: TriCore EABI information regarding how function is called, register allocation for formal parameters and return parameters.

- **TriCore.td**: Contains information about TriCore machine.

4.2 Code Generation Process

4.2.1 Instruction selection

During the instruction selection process an in-memory IR is converted into a TriCore selectionDAG. The LLVM intermediate representation is converted in a tree of DAGs (Direct Acyclic Graphs). An example of selection DAG tree is presented in figure 4.5.

Each DAG graph represent a particular function inside the LLVM IR, while a node in a DAG graph represent a single instruction. A wealth of information is attached with each node. The black arrow show the data-dependency between instructions. The blue arrow (called chain) represent the order in which the blocks are scheduled, i.e. which block will execute first and so on. The red arrow is called
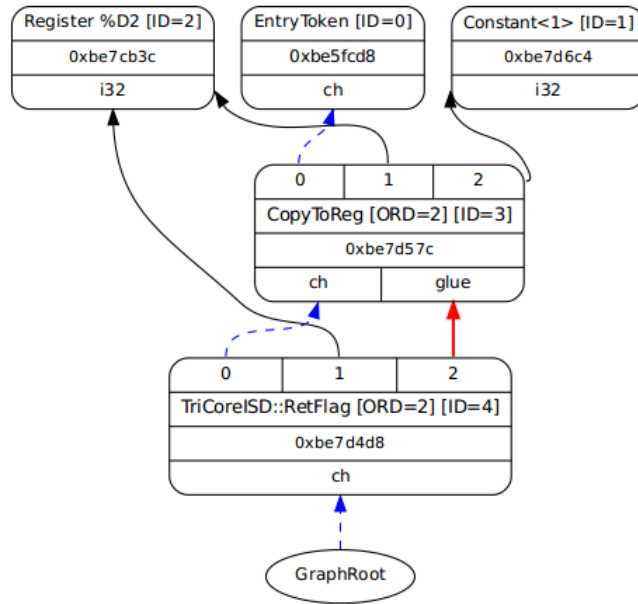
Figure 4.5: instruction selection DAG graph

glue and represents that two nodes are glued together, i.e. two nodes must always be scheduled in the respective aforementioned order, and no other nodes can be scheduled between them. As can be seen from figure 4.5, LLVM tries to add as much target-specific information as possible, but still the SelectionDAG remains target-agnostic for most of the part. This step allows LLVM to apply the tree-based instruction selection algorithm during the instruction selection stage.

4.2.2 Scheduling

The SelectionDAG does not really imply any optimized ordering between different DAG nodes. The first stage of instruction scheduling is called pre-register allocation scheduling. Instruction ordering is performed in this phase with the idea of maximising instruction parallelism. This step help in reducing the total register usage or the number of instruction that are generated in the the first phase. Scheduling happens by reducing the number of live registers to as low as possible. Consider the following IR:

```
%0 = add i32 1, i32 0
%1 = add i32 2, i32 0
store i32 %0, i32* %2, align 4
store i32 %b, i32* %2, align 4
```

Listing 4.5: without instruction re-ordering

In the above example, instruction re-order is yet not performed, hence virtual registers %0, %1, and %2 are live till the end of this basic block. LLVM must

allocate 3 physical registers to each virtual register. Consider the same example with instruction re-ordering:

```
%0 = add i32 1, i32 0
store i32 %0, i32* %2, align 4
%1 = add i32 2, i32 0
store i32 %b, i32* %2, align 4
```

Listing 4.6: with instruction re-ordering

In the listing above, %0 is marked "dead" after the execution of the second instruction. This mean that the same physical register that LLVM allocated in the first instruction can be re-allocated for virtual register %1. As a result, only 2 physical register are used in the process.

Another example of pre-register allocation scheduling is the minimization of TriCore instructions. For example, TriCore's MADD instruction is a combination of a multiply and addition. A MADD instruction can reduce two intructions into a single instruction, and hence reduc total number of clock cycles.

4.2.3   SSA-based Machine Code Optimization

This is an optimization phase that is used to increase the efficiency of the generated code. SSA enforces that each register is only allocated once. Some example of SSA-based optimizations are constant folding, constant sub-expression elimination, constant propagation, and dead code elimination.

4.2.4   Register allocation

LLVM can allocate an infinite number of virtual register in its intermediate representation. These register are to be mapped accordingly to TriCore specific registers. This happens during this stage. LLVM tries to map all pointers to address registers, and data types to data registers. In case, register allocation is not possible, a spill is generated and reported.

Moreover, the second part of the scheduling, called post-register allocation scheduling happens here. As complete register allocation information is available and all virtual registers are removed, instruction re-ordering can be improved further.

4.2.5   Prologue/Epilogue insertion

Prologue and epilogue are similar to headers and footers for a function call. These are added to functions, in order for proper stack allocation. Details about this stage

is given further in this report.

4.2.6 Late machine code optimization

Some final optimizations are performed in this stage. This include peephole optimization. Consider an example where a multiplication is to be performed by 4. One way of accomplishing this is to directly use the *mul* instruction, but such an instruction is really slow, hence reduces code efficiency. A much more efficient way would be to perform a 2 bit right shift. This would allow a further performance increase. Note that, as register allocation is already performed, peephole optimization can only occur if the alternate instruction uses the same number and type of registers.

4.2.7 Code emission

This stage finally emits the TriCore specific assembly files. At this stage LLVM converts the *MachineInstr* into *MCInst* instances. This new form is most suitable for either printing out the assembly file (.s) or object file (.o). All consideration were kept in writing the correct opcodes for each instruction in TriCore, the object file generated in this thesis is ready to be linked together with other object files.

4.3 Target Machine Description

4.3.1 DataLayout String

The DataLayout String provides the overall information about memory layout, data type alignment, and size of pointer in TriCore. It also defines the endianess of an architecture.

```
target datalayout =
    "e-m:e-p:32:32-i1:8:32-i8:8:32-i16:16:32-i64:32-f32:32-f64:32-a:0:32-n32"
```

A small "e" at the start defines that the architecture is little endian. Next it defines the size of pointer that is 32 bit. Later on, is the description of different data types in TriCore, e.g. a char (i8) has a size of 8 bits and an alignment of 32 bits as mentioned in chapter 2. At the end of the string, "n" defines the alignment requirement for the stack that is 32 bits.

4.4 Instruction Selection Details

The

Bibliography

[1] *Tricore architecture volume1: Core architecture v1.3 & v1.3.1*, Infineon Technologies: TriCore Design Group, Bristol, UK, Jan. 2008.

[2] *Tricore architecture manual*, Infineon Technologies: TriCore Design Group, Bristol, UK, Mar. 2007.

[3] (May 2016). The llvm compiler infrastructure, [Online]. Available: `http://llvm.org/`.

[4] (May 2016). Tablegen fundamentals, [Online]. Available: `http://llvm.org/docs/TableGen/index.html`.

[5] C. Erhardt, *Design and Implementation of a TriCore Backend for the LLVM Compiler Framework.*